

PRATIK PAWAR

Senior Data Engineer

pratik.pawar959526@gmail.com | +91 9272469578 | Pune, Maharashtra | linkedin.com/in/prateekpawar-/

PROFESSIONAL SUMMARY

Results-driven Senior Data Engineer with 4+ years of experience designing and delivering scalable ETL/ELT pipelines, cloud data lakes, and enterprise data warehousing solutions. Proven expertise in AWS (S3, Glue, Athena, Lambda), PySpark, Snowflake, and Star Schema modeling. Experienced in building near-real-time ingestion systems using Snowpipe and SQS, and optimizing query performance via Parquet partitioning. Adept at cross-functional collaboration, data quality governance, and KPI reporting for enterprise clients.

CORE COMPETENCIES & SKILLS

Cloud & Big Data: AWS S3, AWS Glue, Amazon Athena, AWS Lambda, IAM | PySpark, Apache Spark, Hadoop (HDFS)

Data Warehousing: Snowflake, Oracle | Snowpipe, COPY INTO, SQS Notifications

Programming & Querying: Python, SQL (MySQL), Spark SQL, PL/SQL

Data Modeling & Engineering: Star Schema, ETL/ELT Pipelines, Data Ingestion, Data Transformation, Data Validation, Parquet Optimization, Data Lineage

PROFESSIONAL EXPERIENCE

Senior Data Engineer

Coforge Limited | Jun 2024 – Present | Pune, India

Enterprise Data Platform — PMI Product Lifecycle & Launch Tracking

- Architected and deployed an enterprise-grade AWS data lake integrating SAP, PLM, E2Open, and Excel sources into Amazon S3, enabling unified product lifecycle tracking for PMI.
- Built data transformation pipelines using AWS Glue (PySpark) for multi-source data cleaning, filtering, and joining, reducing data processing time by improving pipeline modularity.
- Optimized storage costs and query performance by converting raw data into partitioned Parquet format; applied predicate pushdown and partition pruning strategies.
- Loaded curated datasets into Snowflake via COPY INTO for efficient, auditable batch processing with full lineage traceability.
- Designed Star Schema data models (fact and dimension tables) supporting KPI reporting and business intelligence dashboards.
- Conducted data validation, quality checks, and lineage analysis to ensure reporting accuracy and data reliability for enterprise stakeholders.

Key Technologies: AWS S3, AWS Glue, PySpark, Snowflake, Parquet, IAM, Star Schema, Python, SQL

Data Engineer

Vim Soft Private Limited | Feb 2022 – May 2024 | Pune, India

Dynamic Retail Data Insight Platform

- Developed end-to-end ETL pipeline to ingest and process JSON data from client retail systems into Snowflake, enabling structured analytics for business stakeholders.
- Ingested raw data into Amazon S3 and performed SQL-based transformation using AWS Athena, ensuring scalable serverless query execution.
- Performed comprehensive data cleaning including duplicate removal, null handling, and schema standardization to ensure data integrity.
- Implemented Snowpipe with SQS event notifications for near-real-time data ingestion, reducing data latency significantly.
- Optimized Snowflake queries for reporting and analytics, resulting in improved dashboard load times and reduced compute costs.
- Collaborated with business analysts to deliver structured, validated datasets for retail performance and sales insights.

Key Technologies: Snowflake, AWS S3, AWS Athena, Snowpipe, SQS, Python, SQL, ETL Pipelines

CERTIFICATIONS

- IBM Introduction to Data Engineering — IBM / Coursera
- Databricks Lakehouse Fundamentals — Databricks

EDUCATION

Bachelor of Engineering (B.E.) — Civil Engineering

Nashik District Maratha Vidya Prasarak Samaj's College of Engineering, Savitribai Phule Pune University

| 2014 – 2019 | *First Class*

LEADERSHIP & EXTRACURRICULAR

- General Secretary, NDMVP's College of Engineering, Nashik (Jun 2017 – Jun 2018) — Led student initiatives, event management, and cross-department coordination.

ADDITIONAL INFORMATION

Notice Period: 15 Days or Less **Preferred Locations:** Mumbai, Pune, Bengaluru **Languages:** English, Hindi, Marathi (Expert)

Pratik Pawar — Role-Specific Competencies

Position: Associate AI/ML Engineer Product AI & Generative AI at GOLDS Private Limited

Technical Fit Summary:

Based on the core tech-stack requirements for this position, the following matrix outlines the direct alignment of my engineering experience with your team's specifications.

Core Competency / Skill	Hands-on Exp	Relevance & Application Areas
Data Engineer	4.6 Years	Applied extensively in building distributed data architectures, data engineering pipelines, and scale deployments.
AWS	4.6 Years	Applied extensively in building distributed data architectures, data engineering pipelines, and scale deployments.
PySpark	4.6 Years	Applied extensively in building distributed data architectures, data engineering pipelines, and scale deployments.
SQL	4.6 Years	Applied extensively in building distributed data architectures, data engineering pipelines, and scale deployments.
Python	4.6 Years	Applied extensively in building distributed data architectures, data engineering pipelines, and scale deployments.
Spark	4.6 Years	Applied extensively in building distributed data architectures, data engineering pipelines, and scale deployments.
ETL	4.6 Years	Applied extensively in building distributed data architectures, data engineering pipelines, and scale deployments.
Data Pipeline	4.6 Years	Applied extensively in building distributed data architectures, data engineering pipelines, and scale deployments.
Glue	4.6 Years	Applied extensively in building distributed data architectures, data engineering pipelines, and scale deployments.
Redshift	4.6 Years	Applied extensively in building distributed data architectures, data engineering pipelines, and scale deployments.
S3	4.6 Years	Applied extensively in building distributed data architectures, data engineering pipelines, and scale deployments.
Airflow	4.6 Years	Applied extensively in building distributed data architectures, data engineering pipelines, and scale deployments.
Kafka	4.6 Years	Applied extensively in building distributed data architectures, data engineering pipelines, and scale deployments.
Cloud	4.6 Years	Applied extensively in building distributed data architectures, data engineering pipelines, and scale deployments.